

# Advanced Docker Course

---

We are in a containerized world.

Everything is separated.

How do we share/access data ?

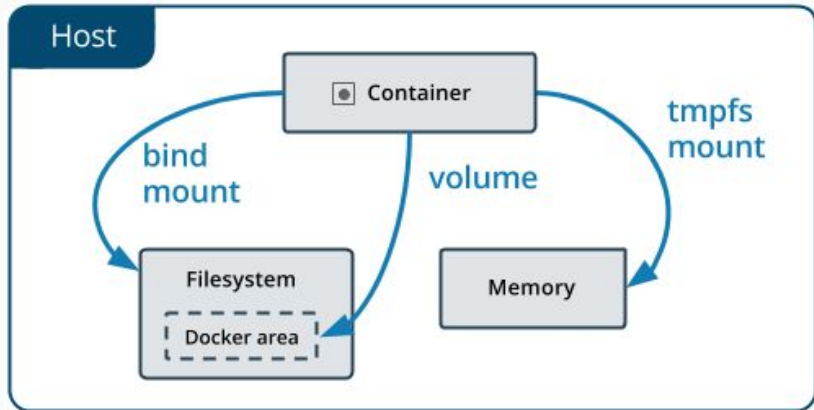
## Data in Docker: Goals

1. How Docker manages data ?
2. Use different type of data storage

Docker has 3 options

- Volumes
- Bind Mounts
- tmpfs mount

# Data in Docker



**Figure 1:** Docker & Data

## Data in Docker: Volumes

- Managed directly by Docker
- Saved in `/var/lib/docker/volumes/` on the host
- System processes can not access data outside Docker
- Recommended by Docker to store data in Docker

## Data in Docker: Volumes

- User can create its own volumes or instruct Docker to create them when required by the containers or services
- The container see the volume as a directory and the user define the name
- Docker guarantees isolation from the host machine
- A volume can be mounted by many containers at the same time

## Data in Docker: Volumes

- If not used it is not destroyed, the user must remove the volume explicitly
- The user can name a volume or Docker create a random name automatically
- A volume can be an *object* in the cloud, the user must use a proper *driver*
- Avoid to increase the size of the container



## Data in Docker: Bind mounts

- Managed directly by the host OS
- Saved in `/in/your/path/` on the host OS
- System processes or Docker containers can access to the data
- It is possible to override important files or directory on the host OS

## Data in Docker: Bind mounts

- Files and directory from the host are mounted inside the container at runtime
- Require the target's full path on the host machine
- On demand creation inside the container
- Very performant
- Rely on the host filesystem
- From inside the container the user has full access to the filesystem: read/write
- Read/Write from outside the Docker container

Expose a *Volume* or *Bind Mount* into the container

```
---volume | -v
```

Docker infers if it is a *Volume* or a *Bind Mount* from the command line

In memory storage

```
--- tmpfs
```

This is useful for ephemeral storage

## Data in Docker: When use What

Volume	Bind mount
Sharing data among containers	Sharing configuration to containers
The host lacks a directory structure	Sharing source code and build products
Data can not be stored locally (cloud)	Stable directory and file structures shared w/ containers

```
$ docker volume
```

Commands:

create     Create a volume

inspect    Display detailed information on one or more  
           volumes

ls         List volumes

prune      Remove all unused volumes

rm         Remove one or more volumes

## Data in Docker: Create a volume

Elixir Docker Advanced Course 2018 Storage

```
$ docker volume create edac18-storage
```

```
edac18-storage
```

## Data in Docker: List volumes

```
$ docker volume ls
```

DRIVER	VOLUME NAME
local	edac18-storage



# Data in Docker: Inspecting a volume

```
$ docker volume inspect edac18-storage

[{"CreatedAt": "2018-09-20T18:38+00:00",
  "Driver": "local",
  "Labels": {},
  "Mountpoint": "/var/lib/docker/volumes/edac18-storage/_data",
  "Name": "edac18-storage",
  "Options": {},
  "Scope": "local"
}]
```

## Data in Docker: Use a volume

Run a container with Ubuntu 18.04 and create a file with something inside

```
$ docker run --rm -v edac18-storage:/data \  
    -it ubuntu:18.04 /bin/bash  
$ echo $RANDOM > /data/seed  
$ exit
```

After closing the container data are stored and can be accessed by another container

```
$ docker run --rm -v edac18-storage:/data \  
    -it ubuntu:18.04 /bin/bash \  
    -c "cat /data/seed"
```

## Data in Docker: Mount points

Load your own directory inside the container

```
$ docker run -v /opt:/host/opt \  
    --name edac18 \  
    --rm -it ubuntu:18.04 /bin/bash
```

## Data in Docker: Mount points

Load your own directory inside the container

```
$ docker run -v /opt:/host/opt \  
    --name edac18 \  
    --rm -it ubuntu:18.04 /bin/bash
```

*Docker does not like relative path*

## Data in Docker: Volumes & Mount points

Combine volumes and mount points in a single instance

```
$ docker run -v /opt:/host/opt \  
             -v edac18-storage:/data \  
             --name edac18 \  
             --rm -it ubuntu:18.04 /bin/bash
```

## Data in Docker: Share data w/ containers

```
$ docker run -v /opt:/host/opt \  
             -v edac18-storage:/data \  
             --name edac18 \  
             --rm -it ubuntu:18.04 /bin/bash
```

Another container can access to the data at the same time

```
$ docker run --volumes-from edac18 \  
             --name backup \  
             --rm -it ubuntu:18.04 /bin/bash
```

*You may notice some lag in updating data, it depends on the underlying Docker filesystem*

## Data in Docker: Share data w/ containers

```
$ docker run -v /opt:/host/opt \  
    -v edac18-storage:/data \  
    --name edac18 \  
    --rm -it ubuntu:18.04 /bin/bash
```

Another container can access to the data and perform a backup automatically

```
$ docker run --volumes-from edac18 \  
    --name backup \  
    --rm -it ubuntu:18.04 \  
    tar vcz /host/opt/backup.tar.gz /data
```

*You may notice some lag in updating data, it depends on the underlying Docker filesystem*

## Data in Docker: Volumes on the fly

When needed you can even create a Volume at runtime w/o using the explicit create command.

```
$ docker run -v bioinfo:/reads  
              -v /opt:/host/opt  
              -v edac18-storage:/data  
              --name edac18  
              --rm -it ubuntu:18.04 /bin/bash
```



## Data in Docker: Volumes on the fly

When needed you can even create a Volume at runtime w/o using the explicit create command.

```
$ docker volume ls
```

DRIVER	VOLUME NAME
local	edac18-storage
local	bioinfo

## Data in Docker: Anonymous Volumes

When needed you can even create a Volume at runtime w/o using the explicit create command.

```
$ docker run -v /anonymous  
             -v bioinfo:/reads  
             -v /opt:/host/opt  
             -v edac18-storage:/data  
             --name edac18  
             --rm -it ubuntu:18.04 /bin/bash
```

## Data in Docker: Anonymous Volumes

When needed you can even create a Volume at runtime w/o using the explicit create command.

```
$ docker volume ls
```

DRIVER	VOLUME NAME
local	edac18-storage
local	bioinfo

*Once exited from the container, there is no evidence about the anonymous container*

## Data in Docker: Anonymous Volumes w/o `--rm`

When needed you can even create a Volume at runtime w/o using the explicit create command.

```
$ docker run -v /anonymous  
             -v bioinfo:/reads  
             -v /opt:/host/opt  
             -v edac18-storage:/data  
             --name edac18  
             -it ubuntu:18.04 /bin/bash
```

## Data in Docker: Anonymous Volumes w/o `--rm`

When needed you can even create a Volume at runtime w/o using the explicit create command.

```
$ docker volume ls
```

DRIVER	VOLUME NAME
local	edac18-storage
local	bioinfo
local	8a76g20jc0gcbgf3952gbdnhd253r801skala7898y...

Docker *volumes*, *images*, *containers* are capped to a default value which is configured at daemon level usually around 10 GB

```
dm.basesize
```

The `dm.basesize` can be increased but the Docker's daemon must be restarted.

The user can run the daemon by hand

```
$ sudo dockerd --storage-opt dm.basesize=50G
```



The user can increase the size but never decrease it.

If the user specifies a value which is lower than the minimum size of *volumes*, *images*, *containers* Docker will complain with errors.

Usually everything works fine, but may happen that Docker's system dir should be wiped out

```
$ sudo service docker stop
```

```
$ sudo rm -rf /var/lib/docker
```