

# Dockerfile

---

## Dockefile: why?

A docker can be created by hand

- pull an image
- start a container
- modify the container
- commit the changes

this is more or less the process for creating a reusable container *an Image*

# Dockerfile: why?

By *hand* is good for practicing or testing but is very bad for

- reproducibility
- automation
- dependencies

# Dockerfile: why?

By *hand* is good for practicing or testing but is very bad for

**reproducibility**: lost the history of the commands that create the final image

- automation
- dependencies

## Dockerfile: why?

By *hand* is good for practicing or testing but is very bad for

- reproducibility

**automation:** images are lost because some disaster and everything was on a local machine

- dependencies

# Dockerfile: why?

By *hand* is good for practicing or testing but is very bad for

- reproducibility
- automation

**dependencies:** images are build from other images and something must be changed in the original image

## Dockerfile: what is it?

A simple text files with all the instructions for

- start (PULL) from a (public) Linux distribution
- install software
- configure the installation
- configure the container for running automatically

# Dockerfile: what is it?

The most minimal Dockerfile

```
FROM ubuntu:18.04
```

*Note* by convention Dockerfile is the name to use for the file containing the instructions.



# Dockerfile: how to use it

A Dockerfile can be used for building a new image

```
$ docker build -t origin .  
Sending build context to Docker daemon 2.048kB  
Step 1/1 : FROM ubuntu:18.04  
--> cd6d8154f1e1  
Successfully built cd6d8154f1e1  
Successfully tagged origin:latest
```

A new image is created with the sha256 cd6d8154f1e1 called origin and the version, in this case by default Docker assign the tag latest

# Dockerfile: FROM

Start from a Linux distribution or previous *installations*/images

```
FROM <image> [AS <name>]
```

```
FROM <image>[:<tag>] [AS <name>]
```

```
FROM <image>[@<digest>] [AS <name>]
```

## Dockerfile: LABEL

Metadata are useful in order to describe the image, making it more consumable by others

`LABEL <key>=<value> <key>=<value> <key>=<value> ...`

```
LABEL org.ingm.group="Your Boss name"
LABEL maintainer="Bonnal Raoul J.P. <bonnal@ingm.org>"
LABEL project="Elixir Test"
LABEL description="Dockerfile example \
with multiple lines."
LABEL version="1.2"

LABEL maintainer"=bonnal@ingm."org
```

User is free to use any kind of key=val convention but the *reverse DNS* notation.

## Dockerfile: ENV

Environment variables can be set inside the container

ENV <key> <value> ENV <key>=<value> ...

```
ENV software="samtools" description=A\ great\ piece\ of\ softw  
author=someone
```

and

```
ENV software samtools  
ENV description A great piece of software  
ENV author someone
```

these variable are available during the building process and when the container is running

## Dockerfile: WORKDIR

Sets the working directory for the following *instructions*

```
ENV MYSUBDIR mytmp  
RUN mkdir /opt/$MYSUBDIR  
WORKDIR /opt/$MYSUBDIR  
RUN pwd
```

Works for RUN, CMD, ENTRYPOINT, COPY and ADD

## Dockerfile: injecting files

To fully customize the image, external files can be included. To achieve this Docker provides two different tools

- ADD
- COPY

## Dockerfile: ADD

```
ADD [--chown=<user>:<group>] <src>... <dest>  
ADD [--chown=<user>:<group>] ["<src>",... "<dest>"]
```

- Digest URLs, download
- Unpack archives (identity, gzip, bzip2 or xz)
- Does not perform authentication
- At every build it is re executed

```
COPY [--chown=<user>:<group>] <src>... <dest>  
COPY [--chown=<user>:<group>] ["<src>",... "<dest>"]
```

- Relative path outside of context does not work
- Works only with local files or directory
- Can copy files from source location to a previous build stage "FROM"
- NO URLs
- NO auto unpacking



When commands must be run with a different shell

```
SHELL ["executable", "parameters"]
```

Set the USER to use during when the containers run. It also set the user for RUN, CMD, ENTRYPOINT following the declaration of USER

```
USER <user>[:<group>]
```

```
USER <UID>[:<GID>]
```

To customise the installation the user must execute commands.

The commands are run inside a default shell `/bin/sh -c`

Use the `SHELL` clause to change the shell for the following Dockerfile

When a `RUN` succeed Docker will write a layer.

## Dockerfile: RUN

RUN have two forms:

```
RUN apt-get update
```

or use a more explicit form where parameters are passed in a sort of JSON notation

```
RUN ["apt-get", "update"]
```

The JSON form does not create a shell for the command, so variable can not be substituted. To use the shell substitution call the shell first.

## Dockerfile: RUN

A RUN command can span multiple lines

```
RUN apt-get install -y wget git python3.6
```

```
RUN apt-get install -y wget \  
    git \  
    python3.6
```

## Dockerfile: RUN

A RUN command can be made by multiple commands

```
RUN comamnd1 && command2
```

The RUN will pass and create a layer only if it succeed. Otherwise, Docker will report the original error.

Combining commands and spanning the commands on multiple lines helps in readability and building complex configurations.

```
RUN apt-get update &&\  
    apt-get install -y wget
```

## Dockerfile: RUN

Combining commands and spanning the commands on multiple lines helps in readability and building complex configurations.

```
RUN apt-get update &&\
    apt-get install -y wget \
        git \
        python3.6
```



# Dockerfile: ENTRYPOINT

Defines a container that runs as an executable

Forms:

**exec:** preferred

```
ENTRYPOINT ["executable", "param1", "param2"]
```

**shell:**

```
ENTRYPOINT command param1 param2
```

# Dockerfile: CMD

Defines the default behaviour for the container.

Forms:

**exec:** preferred

```
CMD ["executable","param1","param2"]
```

**default parameters to ENTRYPOINT:**

```
CMD ["param1","param2"]
```

**shell:**

```
CMD command param1 param2
```

## Dockerfile: VOLUME

It is possible to embed the volume definition at build time.

Any change, at build time, after the definition will be discarded.

```
VOLUME ["/path", "..."]
```

```
VOLUME /path_a /path_b
```

Volumes are:

- created automatically at run time
- can be shared between containers with `--volumes-from`
- are anonymous at runtime
- can be inspected looking at `/var/lib/docker/volumes`

# Dockerfile: VOLUME

Example of creating a VOLUME

```
FROM ubuntu:18.04
RUN mkdir /opt/elixir-volume
RUN echo "This is a file with a foo text" > /opt/elixir-volume
VOLUME ["/opt/elixir-volume"]
```

## Dockerfile: context

Context defines what is visible to Docker at the build time. Data inside the context are copied in a temporary place where the building process is working. The building process can see only data in that temporary place.

This process of building the context can take a lot of time if files are big and many.

Avoid:

- huge files
- temporary or working file
- backup

in the context.

A lean context means quick build.

A Dockerfile is a text file and Docker keep tracks of changes in the file.

Most of the *instructions* generate a layer.

Changes to the text are invalidaing all the following *instructions* and they will be re-executed.

## Dockerfile: example

```
FROM ubuntu:18.04

LABEL org.ingm.group="Your Boss name"
LABEL maintainer="Bonnal Raoul J.P. <bonnal@ingm.org>"
LABEL project="Elixir Test"
LABEL description="Dockerfile example"
LABEL version="1.2"

RUN apt-get update &&\
    apt-get install -y wget \
                                git \
                                python3.6
```

## Dockerfile: bulding

```
$ docker build -t origin .
```

```
$ docker build -t origin Dockerfile .
```

```
$ docker build -t origin -f /absolute/path/Dockerfile .
```



# Dockerfile: building

```
Sending build context to Docker daemon 2.048kB
Step 1/5 : FROM ubuntu:18.04
--> cd6d8154f1e1
Step 2/5 : RUN mkdir /elixir-volume
--> Running in d03725826863
Removing intermediate container d03725826863
--> 0d1e233c8703
Step 3/5 : RUN echo "First Volume" >> /elixir-volume/README.txt
--> Running in b67fd3dc8911
Removing intermediate container b67fd3dc8911
--> 652996dee96a
Step 4/5 : VOLUME ["/elixir-volume"]
--> Running in a5a9251a282d
Removing intermediate container a5a9251a282d
--> 1ea05afe508e
Step 5/5 : RUN touch /elixir-volume/Write.After.Volume.Declaration
--> Running in 3a639154a462
Removing intermediate container 3a639154a462
--> 6336950b090b
Successfully built 6336950b090b
Successfully tagged origin:latest
```

Figure 1: DockerBuilding